

MS-SQL DBA

25 November 2024 10:08

- Primary key
- Foreign key of a table must n should be primary key of some other table
- Candidate key - one or more columns that could serve as primary key (one of these is set to primary key and another is called alternate key)
- Alrenate key -another candidate key which is not primary key (when user doesn't know primary key then they can search using alternate key -ex:train name and number)
- Unique key
- Composite key
- Seconady key
- Super key
- Database objects
 - Tables,views,indexes,procedures(using create command all are objects)
 - View - create view viewname
 - Why? - if a person wants to see only particular part of data when there are lot of entries
 - Abstraction acquired by view(end user can't see the actual base table)
 - Simple and complex view
 - Inbuilt ,user defined types
 - Index(key+ address of the key)
 - Culstured -sorted form
 - Non-clustered- non sorted form
 - Sequence - automatic generation (random number), auto increment, usually to generate primary key
- Collection of fields - records ; collection of records - file
- Properties
 - Atomicity -transaction is treated as single (if any single manipulation fails whole transaction is rolled back and no updatation in database)
 - Locks - table lock,row lock, page lock
- OLTP - day to day /daily current operational data
- OLAP - historical data -cannot be changed
- Heap - not stored in sorted manner(data is appended not sorted)
- IOT table(index organized table)- sorts the data
- Normalization (to avoid redundancy,update,insert anomolies(kinda errors),ambiguous) -process of organising and structure the data
 - 1NF - each table cell should contain signle value and each record should be unique value
 - 2NF - be in 1NF and single column primary key that does not functionally dependent on any subset of candidate key relation

dividing the table into two tables and relating each other with primary key and foreign key(when we have composite key and non-key attributes depends on only part of composite key)

Student_ID	Course_ID	Instructor_ID	Instructor_Name	Grade
1	101	1001	Dr. Smith	A
2	102	1002	Dr. Brown	B
1	102	1002	Dr. Brown	A
3	101	1001	Dr. Smith	C

To check if the relation is in **2NF**, we need to identify **partial dependencies**:

- The **primary key** is **Student_ID, Course_ID**.
- The **Grade** is fully dependent on the **entire primary key (Student_ID, Course_ID)**. So, there is no partial dependency for the **Grade** attribute.
- However, **Instructor_ID** depends only on **Course_ID**, which is **part** of the composite primary key (**Student_ID, Course_ID**), and not the **entire** key. This is a **partial dependency**.
- **Instructor_Name** depends on **Instructor_ID**, which in turn depends on **Course_ID**. Since **Instructor_ID** is dependent on **Course_ID** (a part of the primary key), this is also a partial dependency.

Thus, the relation is **not in 2NF** because it has partial dependencies.

- 3NF- be in 2NF and when there is transitive dependency (a->b,b->c and a->c) - non-key attribute dependent on another non-key attribute and indirectly on primary key
 - In order to achieve that non-key attribute directly depends on primary key we make a relationship

EmployeeID	Name	Department	DepartmentHead	DepartmentLocation
101	Alice	HR	John Smith	Building A
102	Bob	IT	Mary Johnson	Building B
103	Carol	IT	Mary Johnson	Building B

- **Primary Key**: EmployeeID
- **Non-key attributes**: Name, Department, DepartmentHead, DepartmentLocation

Issue:

- The **DepartmentHead** and **DepartmentLocation** depend on **Department**, not directly on **EmployeeID**. This is a **transitive dependency** because:
 - **Department → DepartmentHead** (i.e., a department has a specific head),
 - **Department → DepartmentLocation** (i.e., a department is located in a specific place),
 - **EmployeeID → Department** (i.e., each employee is assigned to a department).

So, we have **transitive dependencies**: **EmployeeID → Department → DepartmentHead** and **EmployeeID → Department → DepartmentLocation**.

- In the new structure:
 - The first table now contains **Department** as the primary key.
 - The second table uses **EmployeeID** as the primary key, and **Department** becomes a foreign key that references the **Department** table.

Resulting Tables After Applying 3NF:

1. Employee Table:

EmployeeID	Name	Department
101	Alice	HR
102	Bob	IT
103	Carol	IT

2. Department Table:

Department	DepartmentHead	DepartmentLocation
HR	John Smith	Building A
IT	Mary Johnson	Building B

- BCNF - be in 3NF and
 - **BCNF Rule**: A relation is in **BCNF** if, for every functional dependency $A \rightarrow B$, the determinant A is a **candidate key**. This means that no non-prime attribute (an attribute not part of any candidate key) can act as a determinant for any other attribute in the relation.

Student_ID	Course_ID	Instructor_ID	Instructor_Name
1	101	1001	Dr. Smith
2	101	1001	Dr. Smith
3	102	1002	Dr. Brown
4	103	1003	Dr. Clark

5 101 1001 Dr. Smith

- **Student_ID → Course_ID:** **Student_ID** is a candidate key, so this is fine in BCNF.
- **Course_ID → Instructor_ID:** **Course_ID** is **not** a candidate key because it does not uniquely identify the entire row (you would need both **Student_ID** and **Course_ID** to uniquely identify each row in the relation). Since **Course_ID** is not a candidate key and it determines **Instructor_ID**, this violates BCNF.
- **Instructor_ID → Instructor_Name:** **Instructor_ID** is a candidate key for the **Instructor** table, so this functional dependency is fine in BCNF.

Conclusion:

- In **BCNF**, every **determinant** must be a **candidate key**. If a determinant is not a candidate key, then the relation violates BCNF and needs to be decomposed into smaller relations.
- **RAID** -redundant array of independent disks

➤ T-SQL-(enhanced sql which includes loops,varibales,conditionals,stored procedures-execute batch of queries)

- **EXAMPLE** – Create a table for demonstration

```
CREATE TABLE Employees (
    EmployeeID INT PRIMARY KEY,
    FirstName NVARCHAR(50),
    LastName NVARCHAR(50),
    Department NVARCHAR(50),
    Salary DECIMAL(10, 2)
);
```

-- Insert some sample data

```
INSERT INTO Employees (EmployeeID, FirstName, LastName, Department, Salary)
VALUES
(1, 'John', 'Doe', 'HR', 50000),
(2, 'Jane', 'Smith', 'IT', 60000),
(3, 'Bill', 'Johnson', 'Sales', 45000);
```

-- Declare a variable to hold the department name

```
DECLARE @Department NVARCHAR(50);
SET @Department = 'IT';
```

-- Declare a variable to hold the average salary

```
DECLARE @AverageSalary DECIMAL(10, 2);
```

-- Try to calculate the average salary for a department

BEGIN TRY

-- Calculate the average salary for the given department

```
SELECT @AverageSalary = AVG(Salary)
FROM Employees
WHERE Department = @Department;
```

-- If no rows are found, throw an error

IF @AverageSalary IS NULL

BEGIN

THROW 50000, 'No employees found in this department.', 1;

END

-- Output the average salary

```
PRINT 'Average salary in ' + @Department + ' department: ' + CAST(@AverageSalary AS NVARCHAR(50));
```

END TRY

BEGIN CATCH

-- Error handling if something goes wrong

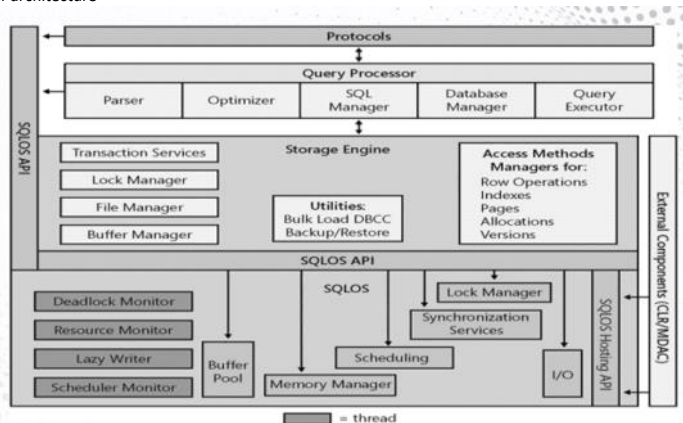
```
PRINT 'Error: ' + ERROR_MESSAGE();
```

END CATCH;

- If there is no value entered '0' is filled unlike sql where null is filled

➤ MSSQL SERVER ARCHITECTURE


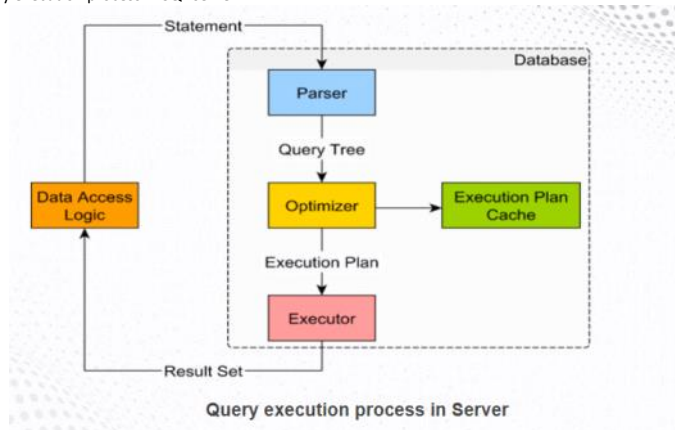
- Client server-based architecture - server response to client query(2 tier architecture-client and database on different systems)
- 1 tier architecture - client and database on same system
- Components of sql server
 - Database engine - creates and drives relational databases
 - Analysis service(ssas)- data analysis componenet
 - Reporting services(ssrs)- generating and deploying report
 - Integration service(ssis) - etl (extract,transform and load)
- Sql server architecture



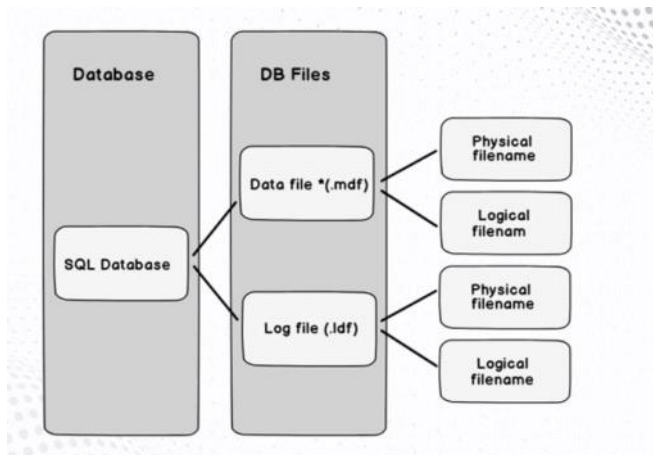
SQL Server Architecture

- Protocol layer-server network interface-(shared memory,tcp/ip,named pipes and tds-tabular data stream:this enables data transfer from client to server machine)
 - Named pipes used by inter-process communication
- Relational engine(cmd parser,optimizer and query executor)
 - Optimizer
 - Phase0- trivial plan
 - Phase1 -transaction processing plan
 - Phase2- parallel processing and optimizations
- Storage engine(file types,access method,buffer manager,plan cache and transaction manager) - responsible for storage and retrieval from storage systems disk/san

- Data file and log file are at disk level
- Buffer manager
 - Plan cache
 - Data parsing
 - Dirty page
- Access method
 - Ddl,select stament - buffer manager
 - Ddl, non-select - transaction manager
- Query execution process in SQL server



- Logical storage
 - The data is organized and accessed (databases ,tables,indexes and schemas)
 - It is how users and applications interact with the data via SQL queries, regardless of the physical storage structure.
 - Data is stored in data file as data pages(8kb)
 - Types of data pages
 - Data page
 - Text/image
 - Index
 - Pfs(page free space)
 - Bcm(bulk change map)
 - Iam(index allocation map)
 - Dcm(differential change map)
 - 1 extent = collection of 8 pages -The extent is the smallest unit of storage that SQL Server allocates to store data on disk.
 - Uniform extent- single object(table) is stored in 8 pages
 - Mixed extent -upto 8 objects can share mixed extent - each of eight pages is owned by different object
 - Global allocation map
 - Shared global allocation map
 - Each page has 96-byte header at the top(free space,page number,page type)
 - Row offset at end of page- stores the base location /postion/address of specific row - stored at end of page in reverse order
- Physical storage
 - It handles the mapping of logical objects (like tables and indexes) to physical storage locations on disk, ensuring data persistence and performance optimization.
 - Data file
 - Physical storage(.mdf)
 - Data files can be created together in file groups fro allocation,distribution and administration purposss
 - Primary data file(.mdf)
 - Data related to objects like tables,indexes,triggres .. Are stored here
 - Every database must have atleast one primary file
 - Secondary data file(.ndf)
 - Optional and contains user-specific data
 - Used to extend the storage of database
 - Log file
 - Stores a series of log records(.ldf)
 - Used for transaction management
 - Transaction log is imp for disaster recovery



- SQL server managers
 - Buffer manager - provides access to the necessary data(soft parse/hard parse)
 - Transaction manager - used when access method concludes request is an operator with no choice
 - Log manager-it monitors all system updates via transaction logs
 - Dirty pages-stored as transaction manager processing logic

➤ MSSQL DATABASE CREATION

- Database types
 - System database
 - Master-
 - Msdb
 - Model
 - Tempdb
 - Resource

Database	Description
Master	This Database contains all of the system-level information for a SQL Server instance.
MSDB	SQL Server Agent used Msdb database for scheduling alerts and job
Model	This database serves as the template for all databases created on the SQL Server instance. Any changes made to the model database, such as database size, collation, recovery model, and other database parameters, are applied to any subsequent databases.
TempDB	TempDB is a workspace for holding users temporary objects or intermediate result sets.
Resource	This Databases is a read-only database that includes system items from SQL Server. System objects are physically stored in the Resource database, but they also appear logically in the sys schema of every database.

- User database
- Add filegroup in sql server
 - Alter database dbname add filegroup fgroup
- Add file to filegroup name

```

ALTER DATABASE Demo
ADD FILE
(
    NAME = demodat3,
    FILENAME = 'C:\Program Files\Microsoft SQL Server
\MSSQL10_50.MSSQLSERVER\MSSQL\DATA\demodat3.ndf',
    SIZE = 5MB,
    MAXSIZE = 100MB,
    FILEGROWTH = 5MB
),
(
    NAME = demodat4,
    FILENAME = 'C:\Program Files\Microsoft SQL Server
\MSSQL10_50.MSSQLSERVER\MSSQL\DATA\demodat4.ndf',
    SIZE = 5MB,
    MAXSIZE = 100MB,
    FILEGROWTH = 5MB
)
TO FILEGROUP demo_secondary;
  
```

- Add log file to sql server

```

ALTER DATABASE Demo
ADD LOG FILE
(
    NAME = Demolog2,
    FILENAME = 'C:\Program Files\Microsoft SQL Server
\MSSQL10_50.MSSQLSERVER\MSSQL\demo_log2.ldf',
    SIZE = 10MB,
    MAXSIZE = 100MB,
    FILEGROWTH = 5%;
)
  
```

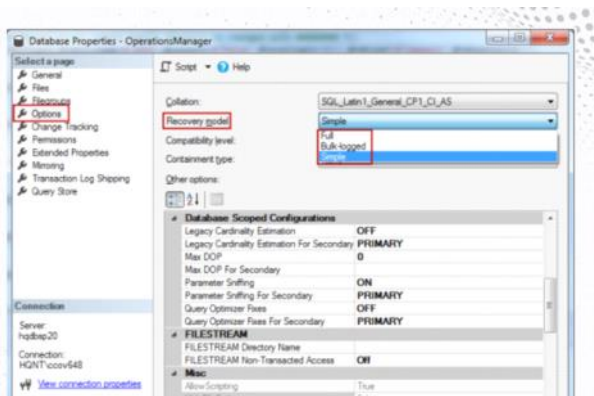
- Modify data file in sql server

```
ALTER DATABASE Demo
MODIFY FILE
(NAME = demodat3,
SIZE = 50MB);
```

```
--Relocating Datafile:

USE master;
ALTER DATABASE demo
MODIFY FILE (name='demo1dat3'
,filename='X:\NewDBFile\demodat3.mdf');
```

- Data operations in sql server
 - Create new db
 - Add new filegroup
 - Add new datafile in primary/secondary file
 - Change datafile size
 - Remove data files
 - Drop database
- MSSQL RECOVERY MODEL
 - Types of recovery model
 - Simple (by default)
 - Full / differential / file backup, no transactional backup
 - Point in time not supported
 - It is used in Developing / testing
 - Full
 - All transactions Ddl + dml documented in transaction log file
 - Supports point in time
 - Data loss can be avoided
 - Bank transactions, financial
 - More filegroups - to recover piecemeal of secondary file / read only file
 - Bulk logged
 - To prevent log file growth
 - Ways to change recovery model
 - Ssms



Steps to take Simple backup in GUI mode

- t-sql

```
SELECT name, recovery_model, recovery_model_desc FROM sys.databases
```

SELECT name, recovery_model, recovery_model_desc FROM sys.databases			
	name	recovery_model	recovery_model_desc
1	master	3	SIMPLE
2	tempdb	3	SIMPLE
3	model	1	FULL
4	msdb	3	SIMPLE
5	OperationsManager	3	SIMPLE
6	OperationsManagerAC	3	SIMPLE
7	OperationsManagerDW	3	SIMPLE

Query to select the recovery model

How to change the recovery model of entire databases of SQL Server instance?

```
EXEC sp_msforeachdb "
IF '?' not in ('tempdb')
begin
    exec ('ALTER DATABASE [?] SET RECOVERY FULL;')
    print '?'
end
"

SELECT name, recovery_model, recovery_model_desc FROM sys.databases
```

- Powershell

```
#load the SqlServer module
Import-Module -Name SqlServer

#input file
$ServerList = "\\hq6021\c$\Server.txt"

#input the recovery model that is intended to search
$RecoveryModel="SIMPLE"

# To exclude system objects
$DatabaseFlag =0

ForEach ($Instance in Get-Content $ServerList)
{
    $s = New-Object ('Microsoft.SqlServer.Management.Smo.Server') $Instance
    $dbs=$s.Databases| Where {$_.RecoveryModel -like "$RecoveryModel" -and $_.IsSystemObject -eq $DatabaseFlag }
    $dbs | select @{Name="ServerName"; Expression = ($Instance)}, Name, RecoveryModel
}


```

SQL Server Command line interface for Change recovery model

```
PS C:\Windows\system32> #load the sqlserver module
Import-Module -Name SqlServer

#input file
$ServerList = "\\hq6021\c$\Server.txt"

#input the recovery model that is intended to search
$RecoveryModel="SIMPLE"

# To exclude system objects
$DatabaseFlag =0

ForEach ($Instance in Get-Content $ServerList)
{
    $s = New-Object ('Microsoft.SqlServer.Management.Smo.Server') $Instance
    $dbs=$s.Databases| Where {$_.RecoveryModel -like "$RecoveryModel" -and $_.IsSystemObject -eq $DatabaseFlag }
    $dbs | select @{Name="ServerName"; Expression = ($Instance)}, Name, RecoveryModel
}


```

ServerName	Name	RecoveryModel
hq6021	Adventureworks2014	Simple
hq6021	Adventureworks2014_datacompress	Simple
hq6021	ApexSQLCrtd	Simple
hq6021	PowerSQL	Simple
hq6021	ReportServerTempdb	Simple
hq6021	operatorksmanger	Simple
hq6021	operatorksmangerAC	Simple
hq6021	operatorksmangerDW	Simple

The changed recovery model for available database

➤ MSSQL Index

- Index will have keys+address
- Clustered index
 - Used to sort and store the data
 - Clustered index keys serves as row finders
- Non-clustered index
 - Row locator is pointer to the row
- Table scan
 - There is system table called sysindexes that contains the information of indexes on tables
 - If no index, then there will be one row which says that there is no index on table
 - When select query is given sql server refers 'indi/index'
- Index creation

- create [unique][clustered /non clustered] index :

```
<indexname> on <object name>(<column list>)
[include(<columnlst>)]
[with fillfactor=<n>]
```

By default an index is non-clustered.

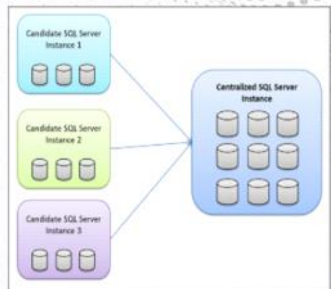
```
create index DNoIndex on Emp(DeptNo)
```

- Single and composite indexes
 - If index is created on single column it is simple index
 - If index is created on conjunction of two tables it is composite table
 - Create index notedidx on emp(deptno asc, job desc)
 - Sp_helpindex emp - to find already created examples
 - Drop index indexname on objectname
- Index used by optimizer
 - Indexes that are well designed can reduce I/O operation and use fewer system resources
 - Query optimizer considers all available data retrieval methods and chooses the best one
 - The approach can be table scan or scan of one or more indexes if exists
 - Query optimizer scans every row in table during a table scan and extracts the row that satisfies the query requirements (it uses a lot of resources and numerous disk IO operations)
 - When query optimizer uses indexes, it searches the index key columns for the storage location of row data

- Index search is better than table search , unlike table index frequently have very few columns per row and rows are sorted
- Monitor index maintenance and update statistics
 - Remove index fragmentation
 - A heavily fragmented index might degrade the performance because it requires additional IO. We can perform reorg or rebuild on index based on its fragmentation values
 - Remove unused indexes
 - An unused index also consumes storage and increases maintenance overhead
 - Below 30% fragmentation we opt for reorganize
 - Above 30% we opt for rebuild
 - Whenever we go for reorganize /rebuild -update statistics is mandatory
 - If update statistics is not performed query optimizer might produce incorrect execution plan
- Index selection and query optimizer
 - Query optimizer uses indexes to produce better execution plan - index selection is the most significant approach in query optimization
 - Unused indexes should be removed
 - Sys.dm_db_index_usage_stats DMV tool: gives the existing unused indexes

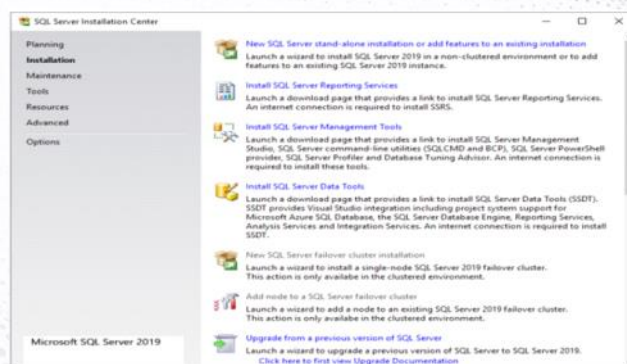
➤ MSSQL Managing Instance

- An instance of the database engine is copy of sqlserver.exe executable that runs as an operating system service (copy of server files, databases and security credentials)
- Each instance manages several system databases and one or more user databases.
- Each computer can run multiple instances of database engines

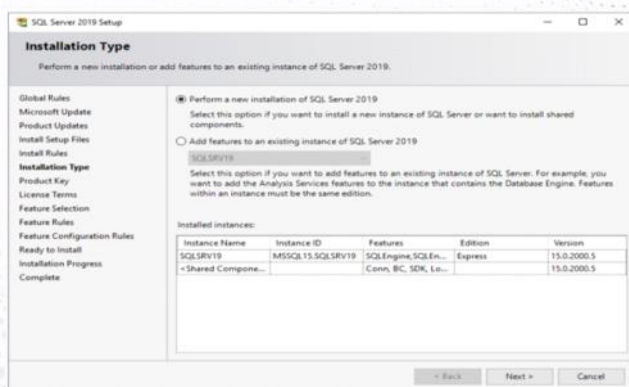


- Multiple instances in sql server
 - Well protected from one another, don't interfere and each instance has own databases ,logins and security
 - Multiple instances gives the ability to spread load over tempdb
- Facts to choose multiple instance
 - Memory usage
 - Tempdb disk usage
 - Cpu usage
 - Io
 - Ease of administration
 - Cost of licensing
 - Monitoring
- Steps to create multiple instances in SQL server

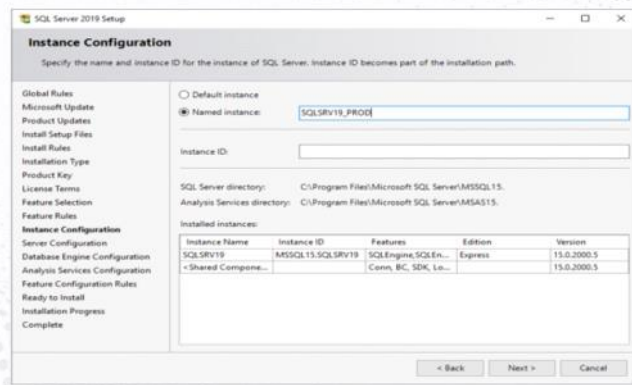
1. Run the Setup file to create multiple instances on the single server.



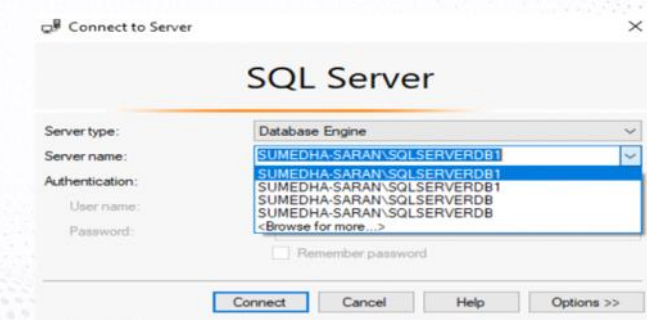
2. Proceed for the new SQL server installation.



3. Select the Named instance option to create new instance.



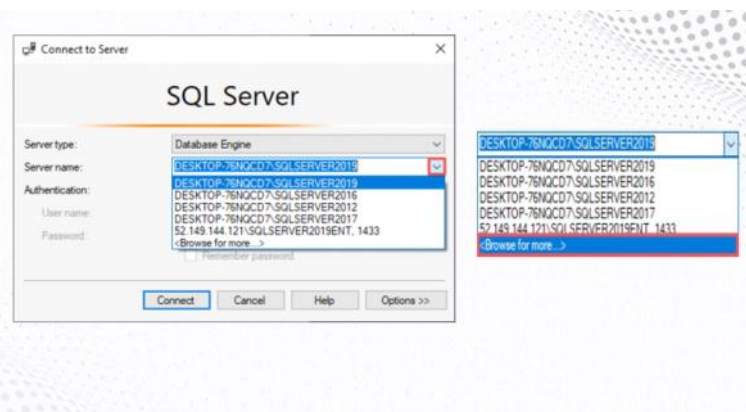
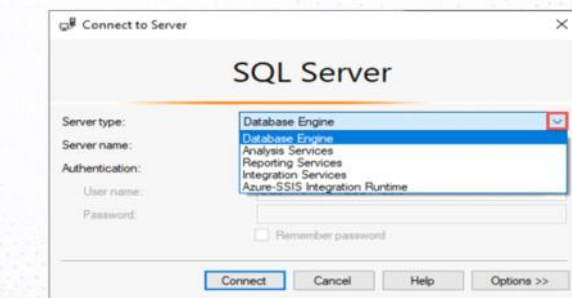
4. Connect with the instance after installation is completed.

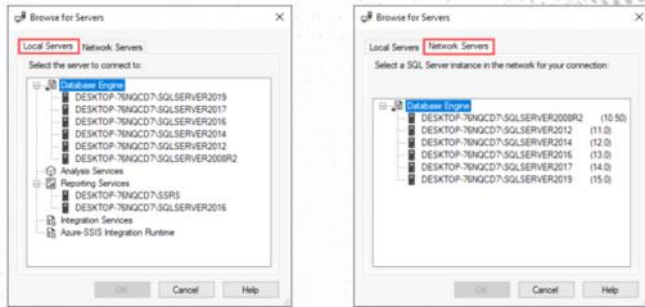


Types of server

- Database engine : core component of sql server instance ,used for storing,processing and securing
- Analysis services - SSAS :is a multidimensional analysis tool that features OLAP , powerful data mining abilities and reporting used in business intelligence to ease your managing data tasks
- Reporting services - SSRS is a reporting and visualization tool for server, and by using it , one can create ,manage and publish reports and dashboards
- Integration services - SSIS is fast and flexible data warehousing tool that can be used for data extraction and executing a wide range of data migration task
- Azure SSIS integration runtime is the compute infrastructure used by azure data factory to provide the data integration capabilities across different network environments

The Following servers available in SQL Server

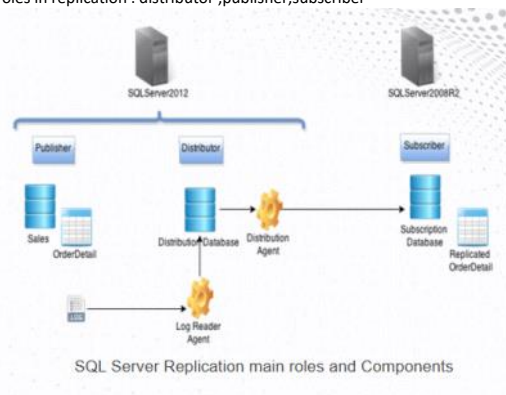




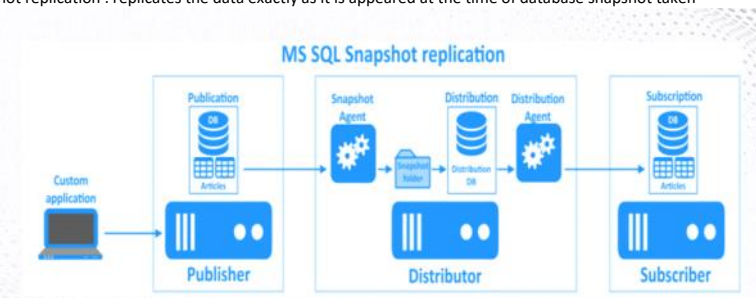
- SQL server components
 - Windows authentication :this mode allows to connect through windows user account
 - SQL server authentication : anytime this mode is used proper login and password is required
 - Azure active directorty
 - Universal with MFA :this mode delivers strong authentication with a different verification options user can choose from - phone call,text message,smart cards with a pin or mobile app notifications
 - Password : use this mode connecting to sql server instance database and when logged into windows using azure AD credentials from a domain not combined with azure
 - Integrated : this mode is used when connecting to sql database and logged into windows using azure active directory credentials from federated domain

➤ MSSQL Replication Overview

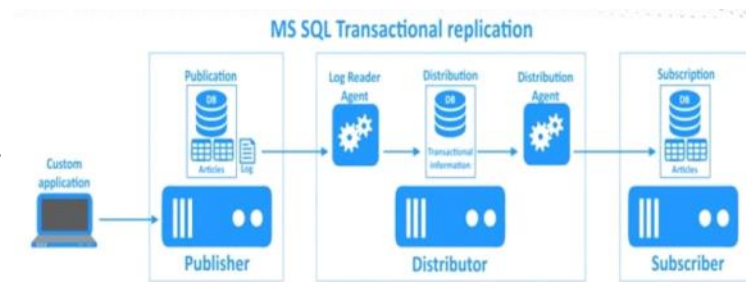
- Replication of group of techniques for transferring data and database objects between databases while synchronizing them to preserve consistency
- Replication is used to distribute data to several locations as well as remote or mobile users via local and wide area networks, dial-up connections, wireless connections, and the internet
- Sql server replication is used for copying and synchronizing data continuously or it can be scheduled to run at predetermined intervals
- There are several different replication techniques that support a variety of data synchronization approches :
 - One-way
 - One-to-many
 - Many-to-one
 - Bi-directional
 - Note :Also keep several datasets in sync with each other
- Transactional replication is commonly applied in server -to-server settings requiring high throughout ,such as:
 - Improving scalability and availability
 - Data warehousing and reporting
 - Integrating data from multiple sites
 - Integrating heterogeneous data
 - Offloading batch processing
- Main roles in replication : distributor ,publisher,subscriber




- Benefits of replication
 - Improves performance and reduces locking conflicts
 - High availability
 - To work with its own copy of data,each location can have own rules and procedures
- Snapshot replication : replicates the data exactly as it is appeared at the time of database snapshot taken



- Advantages :
 - Simple to set up
 - Low maintenance
 - Database tables need not have primary keys
- Disadvt:
 - Modifications made to the subscriber are lost after the new snapshot
 - Snapshot agent generates high impact when it runs
- Transactional replication:
 - Periodic automated replication of data from a master database to a database replica in real time

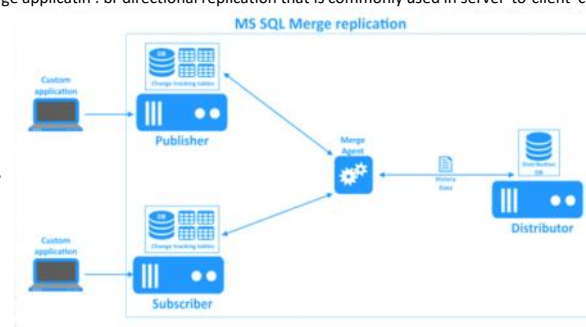


- Advantages:
 - Suitable for high transactional data
 - Requires low latency
- Disadvantages:
 - Setting up transactional replication can be challenging risk, especially for beginners
 - It is prone to faults and can be affected events such as primary key violation
- Peer-to-peer replication: simultaneously replicating the database data to numerous subscribers



Advantage	Disadvantage
Network operating system is not needed	It can slow down the performance for the user because each computer might be being accessed by others.
Expensive server is not needed because individual workstations are used to access the files	Files and folders cannot be centrally backed up
No need for specialist staff such as network technicians because each user sets their own permissions as to which files, they are willing to share.	Files and resources are not organized centrally into a "shared area."
Much easier to set up than a client-server network - does not need specialist knowledge	It is the duty of each individual user to ensure that viruses are not brought into the network.
If one computer fails, it will not disrupt any other part of the network. It just means that those files aren't available to other users at that time.	Apart from the permissions, there is little or no security. Users are frequently not required to log in to their workstations.

- Merge application: bi-directional replication that is commonly used in server-to-client contexts to synchronize data between database servers that are non continuously linked



- Advantages:
 - Let subscribers update records
 - It allows offline subscribers to make change to data
- Disadvantages:
 - Setting up is risk
 - It is prone to create conflicts between publisher and subscribers
- Requirements for sql server replication
 - The following ports should be opened for inbound traffic
 - Tcp 1433,1434,2383,2382,,135,80,443
 - Udp 1434
 - Before installing msmsql server, configure windows firewall and open the relevant poerts for inbound traffic on each host
 - Before setting mssql replication, following softwares should be installed
 - .net framework - a set of libraries
 - Mssql server- the database server software
 - Mssql server management studio (ssms)
- Terminologies in sql

Terminologies	Description
Article	It is the basic unit of SQL Server that contains views, tables, procedures, and so on. It allows you to generate multiple articles on a single object and scale them vertically or horizontally using the filter option.
Publication	A publication is a compilation of articles drawn from a database. It enables you to specify and configure properties at a high level, allowing all articles to easily inherit them.
Publisher Database	The publisher is a database containing a list of objects intended to be replication articles. One or more publications may be available from the publisher. Each publisher develops a number of internal replication procedures that define a data propagation system.
Publisher	It is a database instance that is in charge of distributing data to other sites via replication. It could be storing one or more publications, each of which defines a group of logically connected objects.
Distribution Database	At least one distribution database should be maintained by each distributor. The article information, data, and replication meta-data are all stored in the distribution database. It may contain many distribution databases. All publications from the same publisher, however, must use the same distribution database.

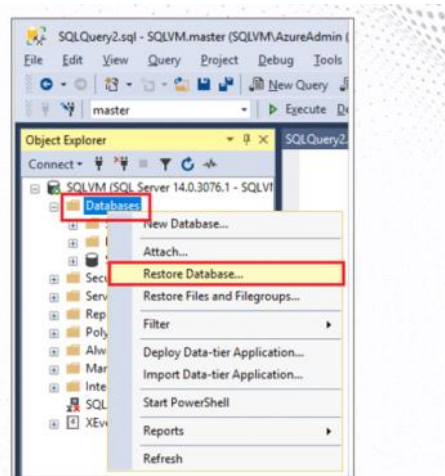
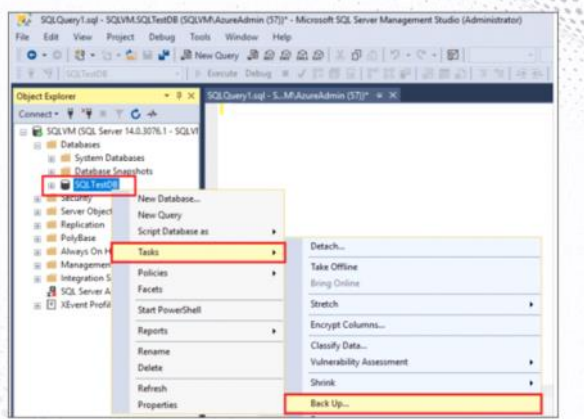
Terminologies	Description
Subscriber	It is a database instance that utilises replication data from a publication. The subscriber can receive data from one or more publications and publishers.
Subscription	It is a request to send a copy of a publication to a subscriber. The subscription specifies when, how, and where the published data will be sent. Subscription database: This is the target database for a replication model

Note: Replication-Details

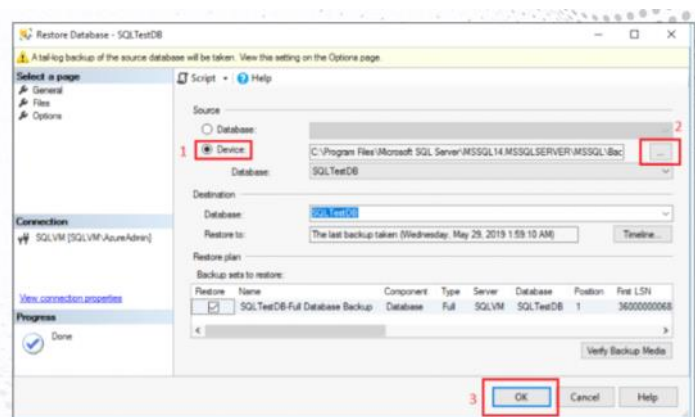
1. Publisher - Instance name
2. Subscriber- Named instance
3. Publication Database- Adventureworks2014
4. Subscription database--Adventureworks2014_Repl
5. Distributor- Instancename
6. Distribution database- Distribution
7. Snapshot folderpath: C:\SQLData\Replidata
8. Tables to be replicated- dbo.sales,humanresource.employee table

➤ MSSQL BACKUP RECOVERY

- Backing up is the process of making duplicates of crucial data on computer .
- Process of restoring deleted or damaged data from backups is called recovery
- Copies of Backups should be stored in secured,off-site place
- Back and recovery strategy should be made and efficiently and tested properly,then only we can implement or use it
- Best practice recommendations
 - Use separate storage
 - Choose appropriate recovery model
 - Design backup strategy
 - Estimate the size of full databse backup
 - Schedule bacjups
 - Test your backups
 - Verify media stability and consistency
 - Document backup/recovery strategy
 - Monitor progress with xevent
- Types of backups
 - Full backups
 - Differential backups
 - File backups
 - Filegroup backups
 - Partila backups
 - Copy-only backups
 - Mirror backups
 - Transaction log backups



Restore the Database from the available backup

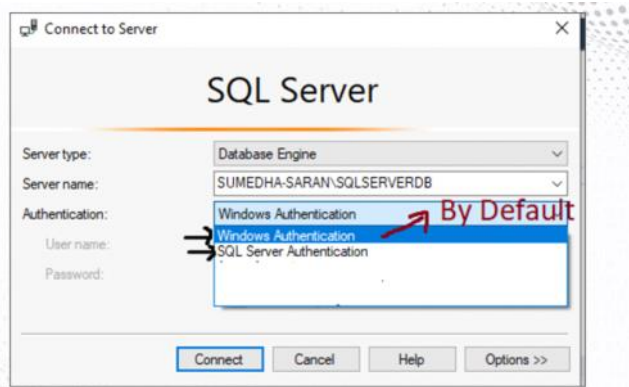


Select the backup for the Database by selecting the device

- Clean resources :
 - Run the following t-sql command to remove the database you created, along with its backup history in the msdb database
 - Exec msdb.dbo.sp_delete_database_backuphistory @database_name='SQL TESTDB'

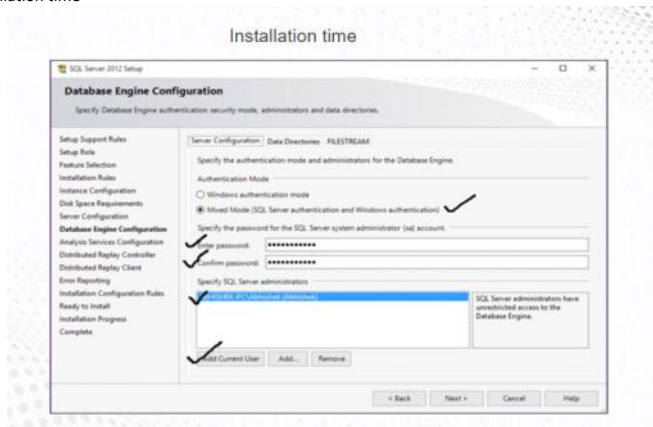
➤ SQL ARCHITECTURE MANAGEMENT

- Sql server architecture management
 - Sql database is created with one or more data files (.mdf and .ndf) and one transaction log file(.ldf). The data files includes the schema and data ,whereas the log file records the most recent modifications
 - Data is organized into pages (8kb)
 - Sql server - read,writes and modifies
- Data retrieval
 - Sql server accesses data by reading a complete page from disc and storing it in memory(RAM)
 - Pages are kept in memory until they are no longer needed
 - As sql works on same data set,same data page is regularly changed or read
- Data modification
 - Sql server uses table based on rows and columns to connect data(data consistency and security)
 - After each db modification - transaction occurrence,data will be saved , and a log record will be written into the transaction log . All the changes will be written sequentially
- Sql server authentication and authorization
 - Supports two authentication modes
 - Windows authentication
 - Sql server authentication

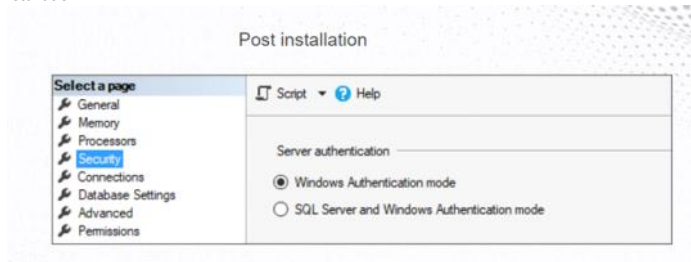


Connecting to SQL Server

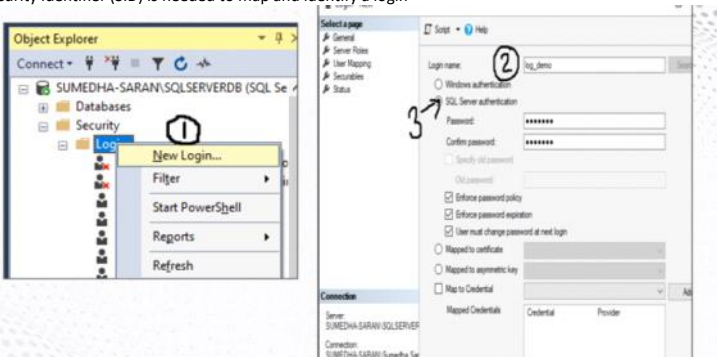
- Windows authentication is integrated with windows user and group accounts, making it possible to use domain or local windows account to log into sql server
- When a windows user connects to sql server ,the db engine validates the login credentials against windows principal authentication, eliminating the need of separate sql credentials
- Setting up authentication:
 - Installation time



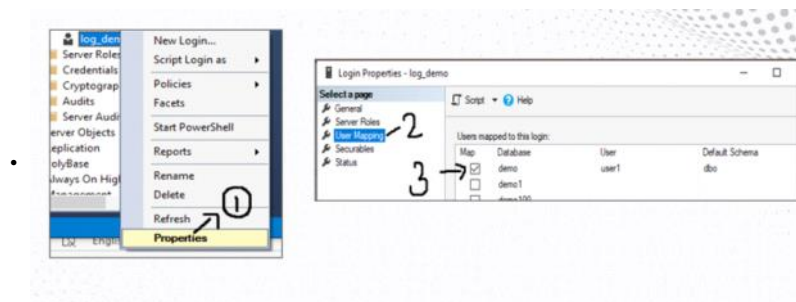
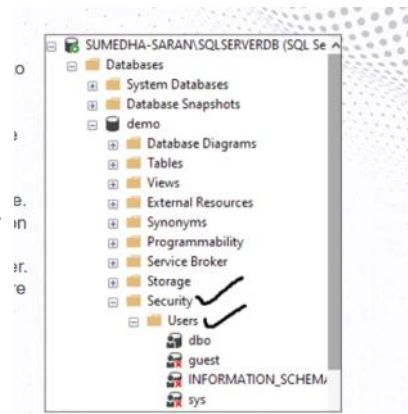
- Post installation



- Sql server login
 - Security identifier (SID) is needed to map and identify a login



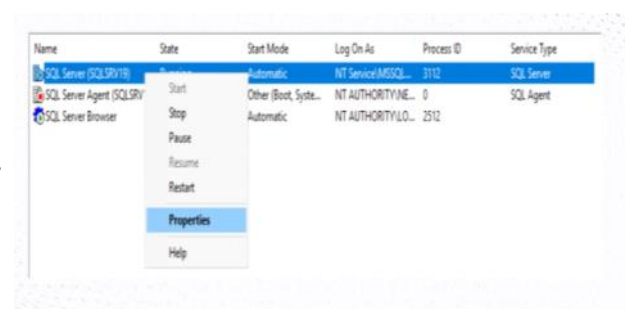
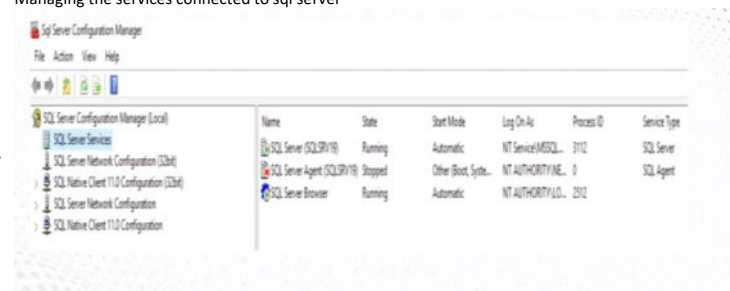
- Server user is different and database user is different



- Sql server permissions
 - Server level permissions
 - It is assigned to login and server roles
 - Database level
 - Is assigned to database user and database roles

➤ MSSQL CONFIGURATION MANAGER


- It is used for
 - Managing network access for sql server client computers
 - Configuring the network protocols used by sql server
 - Managing the services connected to sql server



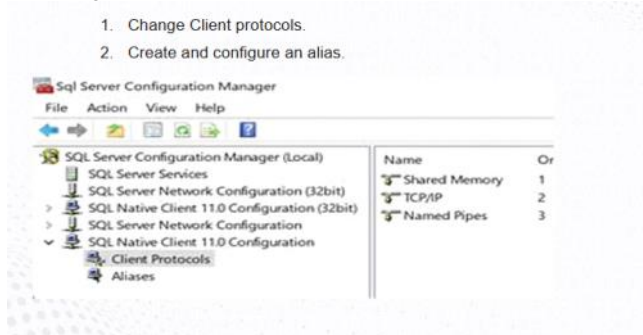
- To change sql server startup parameters



- To enable always on



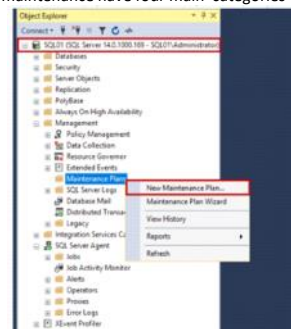
- Native client configuration



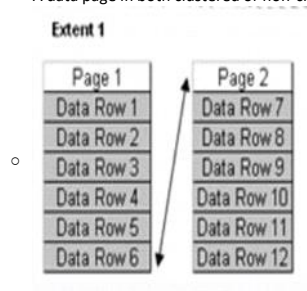
- Clients can connect to the sql server using below
 - Tcp/ip
 - Named pipes
 - ◇ These are used for inter-process communication and is configured for a local area network
 - Shared memory
 - ◇ Clients use the shared memory protocol to connect to sql server instance on the same server
 - ◇ Basic protocol and no configuration options

➤ MSSQL DATABASE MAINTENANCE

- Database maintenance
 - A set of tasks that are all run with the intention to improve your database, these are routines meant to help performance, free up disk space, check for data errors, check for hardware faults, update internal statistics
 - Data maintenance have four main categories




- Index defragmentation
 - A data page in both clustered or non-clustered indexes contains pointers to previous and next page



- Logical fragmentation - the logical order of the pages does not correspond their physical order
- Internal fragmentation - the data pages in the index contain free space. This led to an increase in the number of logical reads during the query execution



- Detecting fragmentation :sys.dm_db_index_physical_stats data management function analyse fragmentation
 - ◇ Avg_page_space_used_in_percent shows the average percentage of data storage space used on the page
 - ◇ Avg_fragmentation_in_percent - external index fragmentation
 - ◇ For tables with clustered indexes fragment_count indicates how many continuous data fragments the index has
- Removing index fragmentation
 - ◇ Alter index all on tablename reorganize
 - ◇ Alter index index_name on tablename rebuild



- Log file maintenance
 - When maintaining a databases's transaction log we'll often want to retrieve information about the log so we can verify its settings or track how much log space is being used
 - Select name,size,max_size,growth,is_percent_growth from sys.database_files where type_desc='LOG'
 - Backing up the log file
 - ◇ DBCC SQLPERF - check log file details
 - ◇ BACKUP LOG employeeDB TO DISK ='E:\ogBackup\employeeDB_log.bak';

The following command can be used to Shrink logfile

```
DBCC SHRINKFILE (EmployeeDB_log, 1);
```

- ```
select * from sys.database_files
```

| file_id | file_name | type_desc | data_space_id | name      | physical_name                                    | state | is_auto_grow | size | max_size | growth | is_read_only |
|---------|-----------|-----------|---------------|-----------|--------------------------------------------------|-------|--------------|------|----------|--------|--------------|
| 1       | WAL       | ROWS      | 1             | master    | C:\Program Files\Microsoft SQL Server\MSSQL15.SQ | 0     | ONLINE       | 655  | -1       | 10     | 0            |
| 2       | WAL       | LOG       | 0             | masterlog | C:\Program Files\Microsoft SQL Server\MSSQL15.SQ | 0     | ONLINE       | 256  | -1       | 10     | 0            |

The following command can be used to add/delete logfile.

```
ALTER DATABASE EmployeeDB
ADD LOG FILE
(
 NAME = EmployeeDB_log2,
 FILENAME = 'C:\SqlData\EmployeeDB2.ldf',
 SIZE = 2MB,
 MAXSIZE = 50MB,
 FILEGROWTH = 10%);
```

- ```
select * from sys.database_files
```

Name	size	max_size	growth	is_percent_growth
EmployeeDB_log	128	268435456	10	1
EmployeeDB_log2	256	6400	10	1

```
ALTER DATABASE EmployeeDB
REMOVE FILE EmployeeDB_log2;
```


- File/data compaction
 - When we work with sql db, the file size both grows and shrinks. Each time the file needs more space, it grows by a certain increment. And as data is added to the db (as a chunk), it is saved into any spot that has enough memory to store that chunk. Unfortunately, the data is not always stored close to other data from the same table. The data itself becomes fragmented, just like the indexes
 - To resolve this, we run what is called a 'data compaction', which reorganizes the file by grouping all associated data together. This not only groups data together, but it can also free up space inside the file as well, which can then be reclaimed by the os as free disk on hard disc
- Integrity check
 - For checking logical and physical integrity of the sql server database /to correct errors run - DBCC CHECKDB
 - Below operations are performed by dbcc checkdb command
 - ◇ DBCC CHECKALOC
 - ◇ DBCC CHECKTABLE
 - ◇ DBCC CHECKCATALOG
 - ◇ DBCC CHECKDB
 - The commands follows below steps
 - ◇ Verify the allocated structures such as GAM,SGAM,PFS,IAM
 - ◇ Verify the table consistency and its all indexes
 - ◇ Finally, verify all the system catalog
- Components of sql server
 - HELP
 - FLUCHAUTHCACHE
 - TRACEOFF

- FREESESSIONCACHE
- TRACEON
- FREESYSTEMCACHE
- CLONEDATABASE

➤ MSSQL DATABASE MONITORING

- Assessing server performance is the main aim of database monitoring
- In order to avoid processes that producing tasks and follow performance trends, effective monitoring entails taking frequent snapshots of current performance
- Server monitor
 - The process of monitoring server's system resources , such as CPU usage ,memory consumption ,I/O ,network ,disc usage, and processes
 - Helps in capacity planning
 - For monitoring server availability and data loss
 - Responsiveness of server
 - To understand the server capacity, user load and speed of the server
 - To detect and prevent any issues that might effect the server proactively
 - Reason to use sql server instance and db monitoring
 - Capacity planning
 - Monitor current resource usage
 - A good sql server monitoring plan can help you stay on top of performance, growth, security.
- Sql server monitoring key areas
 - Disk activity
 - Processor utilization
 - Memory
- Sql server monitoring tools
 - Sql server profiler
 - Sql server management studio reports
 - Activity monitor
 - Sql logs,system monitor
 - Database engine stores procedures
- Window application log
 - Related to windows event viewer
 - System log
 - Security log
 - Application log

The Windows Application Event Log is a comprehensive set of notifications and logs pertaining to important system-wide events including SQL Server.



- DBCC(Database console command)
 - The dbcc commands listed below act on read-only database snapshot that the database creates internally
 - DBCC CHECKALLOC
 - DBCC CHECKCATALOG
 - DBCC CHECKDB
 - DBCC CHECKFILEGROUP
 - DBCC CHECKTABLE

➤ MSSQL MOVING DATA

- Transferring data from one db to another db

SYNTAX:

```
Select * into DestinationDB.dbo.tableName from SourceDB.dbo.SourceTable
```

○ Example:

```
-----
Source Database: Demo
Destination database: Sales
target table: Emp
select * into demo.dbo.emp from sales.dbo.Emp
```

- Open SQL Server Management Studio.
- Right-click on the database name, then select "Tasks" > "Export data..." from the object explorer.



- The SQL Server Import/Export wizard opens; click on "Next" to proceed.



- Data transfer using script
 - Right-click on database and select task
 - Click on generate scripts
 - Select the database you want to script
 - Choose the object types
 - Select the tables
 - Select the output option for the script
 - Change/edit the db name to the name you want to execute the script for